

Agentic Coding Overview

Models, harnesses, access, routing, and benchmarks

Cédric Jung Jozef Janus

30 September 2026

Model	Context	Deployment note
DeepSeek V4 Flash 284B	393k tokens	RTX PRO 6000 Blackwell migration
Qwen 3.6 35B	262k tokens	4 replicas, one A40 each
GLM 5.2 744B Preview	262k tokens	8 AMD MI300X GPUs

Transition

Qwen 3.5 397B reached its announced retirement date on 30 September. DeepSeek V4 Flash is its planned replacement on the Blackwell infrastructure.

AQUEDUCT: what may come next

Release-note candidates are not confirmed deployments

Candidate model	Discussed hardware	Status
Qwen3.8-27B	freed H200 node; possibly 4 instances	Candidate
DeepSeek V4 Flash Vision Exp	RTX PRO 6000 Blackwell	Candidate
GLM-5.3-Flash	RTX PRO 6000 Blackwell	Candidate
DeepSeek V4.1 Flash / GLM-5.3	AMD MI300X	Candidate

Before use

Check the live `/v1/models` endpoint. A release-note mention does not make a model available through AQUEDUCT.

Hardware behind the AQUEDUCT models

Memory and replica counts shape what can be served

Accelerator	Memory / GPU	AQUEDUCT use
NVIDIA A40	48 GB	Qwen 3.6 35B: 4 independent one-GPU replicas
RTX PRO 6000 Blackwell	96 GB	DeepSeek V4 Flash replaces the retiring Qwen 3.5 deployment
NVIDIA H200	141 GB	Initially served DeepSeek; freed for a future deployment
AMD MI300X	192 GB	GLM 5.2 744B Preview: 8 GPUs in one replica

Hardware mapping reflects the public overview plus migration notes; verify the live deployment before relying on it.

Source: [TU Wien AQUEDUCT overview](#); [release notes](#)

1. **Request access** through the TU Wien AQUEDUCT access process.
2. **Wait for approval / allowlisting** before using the service.
3. **Use the issued credentials** in an API-compatible client or coding harness.
4. **Check available models** through the live `/v1/models` endpoint.

Practical point

The public model catalog describes the service; it does not grant an account automatically.

Model: decides

- Understands the task and code
- Proposes edits and tool calls
- Interprets tool results

Examples: GPT, Claude, Qwen, DeepSeek

Harness: executes the loop


- Builds context and exposes tools
- Reads files, edits code, runs tests
- Manages state and permissions

Examples: Codex, Claude Code, OpenCode

User task → harness ↔ model → tested change

Source: [OpenAI](#), *Unrolling the Codex agent loop*


Codex as a concrete example



- Coding post-training can use trajectories: inspect, edit, run tests, observe, revise.
- OpenAI describes codex-1 and GPT-5-Codex as trained or optimized for agentic coding in tool-using environments.
- A benchmark therefore measures the **model, harness, tools, and reasoning budget** together.

Why LLM routing platforms exist

One integration, several model and provider choices



- **Model** = which neural network responds.
- **Provider** = who runs it; providers vary in price, speed, region, and retention rules.
- The router can select a provider, apply constraints, and fail over when needed

What it offers

- One API across many models
- Provider ordering and filtering
- Price or throughput priority
- Provider and model fallback

Routing choices

- Cheapest eligible endpoint
- Higher throughput endpoint
- Specific provider or region
- EU in-region option for eligible enterprise use

Check the chosen provider and region before sending sensitive code or data.

Source: [OpenRouter quickstart](#); [sovereign AI docs](#)

Different ways to constrain where inference runs

Service	Role	Regional control
EUrouter	Multi-provider model gateway	EU processing by design; provider filters
Eden AI	Multi-provider AI gateway	Dedicated European endpoint
Scaleway	Inference provider with many models	European-hosted serving; not a multi-provider router

Selection question

Specify processing location, provider ownership, and retention requirements separately.
"EU-hosted" and "EU-owned" are different constraints.

Source: [EUrouter docs](#); [Eden AI EU](#); [Scaleway FAQ](#)

Subscription: fixed monthly fee

ChatGPT Plus: \$20/month

Codex usage included, subject to plan limits.

API: metered tokens

GPT-5.6 Sol: \$4/M input + \$20/M output

5M input + 1M output = \$40.

For frontier models, a sufficiently active developer can therefore consume **more than \$20 of API tokens very quickly**, while the same \$20 subscription may continue to provide included usage until its rate or weekly limits are reached.

At the moment, subscription pricing can offer dramatically better effective value for a human using an AI coding agent interactively than paying standard API rates for every token.

Source: ChatGPT Plus; Codex plan usage; API pricing

Which benchmarks answer which questions?

One result slide per benchmark follows


Benchmark	Metric	What it measures
GDPval-AA v2.1	Elo	Professional written knowledge work (writing proxy)
DeepSWE v1.1	Resolved %	Long-horizon software engineering
Terminal-Bench 4.0	Pass rate %	Agentic terminal tasks; common evaluator
Terminal-Bench 2.1	Pass rate %	Older version with more open-model coverage
Humanity's Last Exam	pass@1 %	Broad expert academic reasoning
GPQA Diamond	Accuracy %	Graduate-level physics, biology, chemistry

Charts: red =

open-weight comparison models; blue = proprietary models. Compare exact versions; N/A means no public exact-model result, not zero.


Source: [Terminal-Bench](#); [DeepSWE](#); [Artificial Analysis](#)

Professional written knowledge-work proxy; Elo, higher is better



Same benchmark version and evaluator across all eight models. This is not a direct academic-paper-writing test.

Source: [Artificial Analysis](#), GDPval-AA v2.1




Qwen: no public exact-model v1.1 result. Harnesses and reasoning budgets differ across reported scores.

Source: [DeepSWE leaderboard](#); model-card and vendor results cited in supplied benchmark notes

Terminal-Bench 4.0

Agentic terminal tasks; approximate pass rate in Artificial Analysis setup




Qwen: no exact v4.0 score. Scores remain sensitive to harness, tools, and reasoning settings.

Source: [Artificial Analysis, Terminal-Bench 4.0](#)

Terminal-Bench 2.1

Agentic terminal tasks; published results under different setups




Several exact-model scores are unavailable. Do not compare v2.1 percentages with v4.0 percentages.

Source: [Terminal-Bench](#); model-card and vendor results cited in supplied benchmark notes


Humanity's Last Exam

Expert academic reasoning; text-only subset pass@1, higher is better



Artificial Analysis uses a 2,158-question text-only subset. HLE spans humanities and mathematics as well as science.

Source: [Artificial Analysis](#), [Humanity's Last Exam](#); [Qwen model card](#)



198 questions in physics, biology, and chemistry. GPT-5.6 Sol value is the Artificial Analysis result.

Source: [Artificial Analysis](#), GPQA Diamond; Qwen model card

- [TU Wien AQUEDUCT overview and release notes](#)
- [OpenAI: Unrolling the Codex agent loop](#)
- [OpenAI: GPT-5-Codex system-card addendum](#)
- [OpenRouter quickstart and regional routing](#)
- [EUrouter documentation, Eden AI EU, Scaleway FAQ](#)
- [ChatGPT Plus and GPT-5.6 Sol API pricing](#)

Deck compiled from the repository's [documentation](#)/ [research notes](#), dated 30 September 2026. Links are clickable in the PDF.

Exact versions and configurations matter

- Artificial Analysis: GDPval-AA v2.1
- DataCurve: DeepSWE v1.1
- Artificial Analysis: Terminal-Bench 4.0
- Artificial Analysis: Humanity's Last Exam
- Artificial Analysis: GPQA Diamond
- Terminal-Bench project

Interpretation

Agentic scores reflect model, harness, tools, and reasoning budget. Benchmark versions and evaluators must match for a fair comparison.