

Agentic Coding Overview

Models, harnesses, access, routing, and benchmarks

Cédric Jung Jozef Janus

30 September 2026

Model	Context	Deployment note
DeepSeek V4 Flash 284B	393k tokens	RTX PRO 6000 Blackwell migration
Qwen 3.6 35B	262k tokens	4 replicas, one A40 each
GLM 5.2 744B Preview	262k tokens	8 AMD MI300X GPUs

Transition

Qwen 3.5 397B reached its announced retirement date on 30 September. DeepSeek V4 Flash is its planned replacement on the Blackwell infrastructure.

Release-note candidates are not confirmed deployments

Candidate model	Discussed hardware	Status
Qwen3.8-27B	freed H200 node; possibly 4 instances	Candidate
DeepSeek V4 Flash Vision Exp	RTX PRO 6000 Blackwell	Candidate
GLM-5.3-Flash	RTX PRO 6000 Blackwell	Candidate
DeepSeek V4.1 Flash / GLM-5.3	AMD MI300X	Candidate

Before use

Check the live `/v1/models` endpoint. A release-note mention does not make a model available through AQUEDUCT.

Hardware behind the AQUEDUCT models

Memory and replica counts shape what can be served

Accelerator	Memory / GPU	AQUEDUCT use
NVIDIA A40	48 GB	Qwen 3.6 35B: 4 independent one-GPU replicas
RTX PRO 6000 Blackwell	96 GB	DeepSeek V4 Flash replaces the retiring Qwen 3.5 deployment
NVIDIA H200	141 GB	Initially served DeepSeek; freed for a future deployment
AMD MI300X	192 GB	GLM 5.2 744B Preview: 8 GPUs in one replica

Hardware mapping reflects the public overview plus migration notes; verify the live deployment before relying on it.

Source: [TU Wien AQUEDUCT overview](#); [release notes](#)

1. **Request access** through the TU Wien AQUEDUCT access process.
2. **Wait for approval / allowlisting** before using the service.
3. **Use the issued credentials** in an API-compatible client or coding harness.
4. **Check available models** through the live `/v1/models` endpoint.

Practical point

The public model catalog describes the service; it does not grant an account automatically.

Model: decides

- Understands the task and code
- Proposes edits and tool calls
- Interprets tool results

Examples: GPT, Claude, Qwen, DeepSeek

Harness: executes the loop


- Builds context and exposes tools
- Reads files, edits code, runs tests
- Manages state and permissions

Examples: Codex, Claude Code, OpenCode

User task → harness ↔ model → tested change

Source: [OpenAI](#), *Unrolling the Codex agent loop*


Codex as a concrete example



- Coding post-training can use trajectories: inspect, edit, run tests, observe, revise.
- OpenAI describes codex-1 and GPT-5-Codex as trained or optimized for agentic coding in tool-using environments.
- A benchmark therefore measures the **model, harness, tools, and reasoning budget** together.

Why LLM routing platforms exist

One integration, several model and provider choices



- **Model** = which neural network responds.
- **Provider** = who runs it; providers vary in price, speed, region, and retention rules.
- The router can select a provider, apply constraints, and fail over when needed

What it offers

- One API across many models
- Provider ordering and filtering
- Price or throughput priority
- Provider and model fallback

Routing choices

- Cheapest eligible endpoint
- Higher throughput endpoint
- Specific provider or region
- EU in-region option for eligible enterprise use

Check the chosen provider and region before sending sensitive code or data.

Source: [OpenRouter quickstart](#); [sovereign AI docs](#)

Different ways to constrain where inference runs

Service	Role	Regional control
EUrouter	Multi-provider model gateway	EU processing by design; provider filters
Eden AI	Multi-provider AI gateway	Dedicated European endpoint
Scaleway	Inference provider with many models	European-hosted serving; not a multi-provider router

Selection question

Specify processing location, provider ownership, and retention requirements separately.
"EU-hosted" and "EU-owned" are different constraints.

Source: [EUrouter docs](#); [Eden AI EU](#); [Scaleway FAQ](#)

Subscription: fixed monthly fee

ChatGPT Plus: \$20/month

Codex usage is included within plan limits. Good fit for interactive individual use.

API: metered tokens

GPT-5.6 Sol: \$4/M input + \$20/M output

Example: 5M input + 1M output = \$40.
Good fit for programmable workflows.

The subscription is not an API credit balance: it has product and usage limits. Caching and workload shape API cost. Recheck rates before purchasing.

Source: [ChatGPT Plus](#); [GPT-5.6 Sol API pricing](#)

Which benchmarks answer which questions?

Keep versions, metrics, and evaluators visible

Benchmark	Measure	Use / limitation
GDPval-AA v2.1	Elo	Written knowledge work; proxy, not academic paper writing
DeepSWE v1.1	% tasks resolved	Long-horizon software engineering; harness varies
Terminal-Bench 4.0	% tasks passed	Agentic terminal work; common evaluator for chart
Humanity's Last Exam	pass@1 %	Broad expert reasoning; not science-only


Reading rule

Compare exact model revisions under the same benchmark version. Missing results mean *no public exact-model score located*, not zero.

Source: [Terminal-Bench](#); [DeepSWE](#); [GDPval-AA](#)

Agentic coding result: Terminal-Bench 4.0

Approximate task pass rate, Artificial Analysis setup



Red: open-weight comparison models; blue: proprietary frontier models. Qwen3.6-35B-A3B has no exact v4.0 result. Scores are approximate and depend on harness and reasoning settings.

Source: [Artificial Analysis](#), Terminal-Bench 4.0; model-name mapping in supplied benchmark notes

- [TU Wien AQUEDUCT overview and release notes](#)
- [OpenAI: Unrolling the Codex agent loop](#)
- [OpenAI: GPT-5-Codex system-card addendum](#)
- [OpenRouter quickstart and regional routing](#)
- [EUrouter documentation, Eden AI EU, Scaleway FAQ](#)
- [ChatGPT Plus and GPT-5.6 Sol API pricing](#)

Deck compiled from the repository's [documentation](#)/ [research notes](#), dated 30 September 2026. Links are clickable in the PDF.

- Artificial Analysis: GDPval-AA v2.1
- DataCurve: DeepSWE v1.1
- Artificial Analysis: Terminal-Bench 4.0
- Artificial Analysis: Humanity's Last Exam
- Terminal-Bench project

Interpretation

Agentic scores reflect model, harness, tools, and reasoning budget. Benchmark versions and evaluators must match for a fair comparison.